







Embedding table grows as 4^k · backbone stays fixed

Parameter composition per tokenizer resolution · $n_embd = 48$, $n_layer = 2$, $n_head = 2$



Embedding table (vocab_size × 48 params)

Transformer backbone (60,480 params, fixed)